

STEI - Institut Teknologi Bandung

Modul 5

Intel X86-64 Bit Machine

Machine-Level Programming

EL3011 Arsitektur Sistem Komputer




Contents:

1. Intel X86 Evolution <https://www.youtube.com/watch?v=re9g5ZlHeLk>
2. Data Types <https://www.youtube.com/watch?v=crBBaHyib7Q>
3. Registers <https://www.youtube.com/watch?v=hjwzIK7KSp8>
4. Processor Operations <https://www.youtube.com/watch?v=zMS3isOty1I>
5. mov Operations <https://www.youtube.com/watch?v=w884J7U0RRM>
6. Addressing Modes <https://www.youtube.com/watch?v=jS26ecAFZcl>
7. Addressing Computation https://www.youtube.com/watch?v=61B_r-B75ZM
8. Arithmetic and Logical Operations <https://www.youtube.com/watch?v=Uv-mUzL6w3c>
9. C Programming, Assembly and Machine Codes <https://www.youtube.com/watch?v=OmjlUHDxe3I>

This module adopted from 15-213 Introduction to Computer Systems Lecture, Carnegie Mellon University, 2020





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.1. Intel X86 Evolution

EL3011 Arsitektur Sistem Komputer



Intel x86 Processors

- Dominate laptop/desktop/server market
- Evolutionary design
 - Backwards compatible up until 8086, introduced in 1978
 - Added more features as time goes on
- x86 is a Complex Instruction Set Computer (CISC)
 - Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
- Compare: Reduced Instruction Set Computer (RISC)
 - RISC: *very few* instructions, with *very few* modes for each
 - RISC can be quite fast (but Intel still wins on speed!)
 - Current RISC renaissance (e.g., ARM, RISC V), especially for low-power



Intel x86 Evolution: Milestones


<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
• 8086	1978	29K	5-10
<ul style="list-style-type: none">• First 16-bit Intel processor. Basis for IBM PC & DOS• 1MB address space			
• 386	1985	275K	16-33
<ul style="list-style-type: none">• First 32-bit Intel processor , referred to as IA32• Added “flat addressing”, capable of running Unix			
• Pentium 4E	2004	125M	2800-3800
<ul style="list-style-type: none">• First 64-bit Intel x86 processor, referred to as x86-64			
• Core 2	2006	291M	1060-3333
<ul style="list-style-type: none">• First multi-core Intel processor			
• Core i7	2008	731M	1600-4400
<ul style="list-style-type: none">• Four cores			



Intel x86 Processors

- Machine Evolution

		Transistor
• 386	1985	0.3M
• Pentium	1993	3.1M
• Pentium/MMX	1997	4.5M
• PentiumPro	1995	6.5M
• Pentium III	1999	8.2M
• Pentium 4	2000	42M
• Core 2 Duo	2006	291M
• Core i7	2008	731M
• Core i7 Skylake	2015	1.9B
- Added Features
 - Instructions to support multimedia operations
 - Instructions to enable more efficient conditional operations
 - Transition from 32 bits to 64 bits
 - More cores



Intel x86 Processors, cont.

• Past Generations Process technology

- | | | |
|-------------------------------|------|--------|
| • 1 st Pentium Pro | 1995 | 600 nm |
| • 1 st Pentium III | 1999 | 250 nm |
| • 1 st Pentium 4 | 2000 | 180 nm |
| • 1 st Core 2 Duo | 2006 | 65 nm |


Process technology dimension
= width of narrowest wires
(10 nm \approx 100 atoms wide)

• Recent & Upcoming Generations

- | | | |
|-----------------|------|-------|
| 1. Nehalem | 2008 | 45 nm |
| 2. Sandy Bridge | 2011 | 32 nm |
| 3. Ivy Bridge | 2012 | 22 nm |
| 4. Haswell | 2013 | 22 nm |
| 5. Broadwell | 2014 | 14 nm |
| 6. Skylake | 2015 | 14 nm |
| 7. Kaby Lake | 2016 | 14 nm |
| 8. Coffee Lake | 2017 | 14 nm |
| 9. Cannon Lake | 2018 | 10 nm |
| 10. Ice Lake | 2019 | 10 nm |
| 11. Tiger Lake? | 2020 | 10 nm |



2018 State of the Art: Coffee Lake



- **Mobile Model: Core i7**

- 2.2-3.2 GHz
- 45 W

- **Desktop Model: Core i7**

- Integrated graphics
 - 2.4-4.0 GHz
 - 35-95 W

- **Server Model: Xeon E**

- Integrated graphics
 - Multi-socket enabled
 - 3.3-3.8 GHz
 - 80-95 W



x86 Clones: Advanced Micro Devices (AMD)

- Historically
 - AMD has followed just behind Intel
 - A little bit slower, a lot cheaper
- Then
 - Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
 - Built Opteron: tough competitor to Pentium 4
 - Developed x86-64, their own extension to 64 bits
- Recent Years
 - Intel got its act together
 - 1995-2011: Lead semiconductor “fab” in world
 - 2018: #2 largest by \$\$ (#1 is Samsung)
 - 2019: reclaimed #1
 - AMD fell behind
 - Relies on external semiconductor manufacturer Global Foundries
 - ca. 2019 CPUs (e.g., Ryzen) are competitive again




Intel's 64-Bit History

- 2001: Intel Attempts Radical Shift from IA32 to IA64
 - Totally different architecture (Itanium, AKA “Itanic”)
 - Executes IA32 code only as legacy
 - Performance disappointing
- 2003: AMD Steps in with Evolutionary Solution
 - x86-64 (now called “AMD64”)
- Intel Felt Obligated to Focus on IA64
 - Hard to admit mistake or that AMD is better
- 2004: Intel Announces EM64T extension to IA32
 - Extended Memory 64-bit Technology
 - Almost identical to x86-64!
- Virtually all modern x86 processors support x86-64
 - But, lots of code still runs in 32-bit mode



Next segment Data Types





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.2. Data Types

EL3011 Arsitektur Sistem Komputer




Levels of Abstraction


C programmer

```
#include <stdio.h>
int main(){
    int i, n = 10, t1 = 0, t2 = 1, nxt;
    for (i = 1; i <= n; ++i){
        printf("%d, ", t1);
        nxt = t1 + t2;
        t1 = t2;
        t2 = nxt; }
    return 0; }
```


Assembly programmer



Computer designer



Gates, clocks, circuit layout, ...

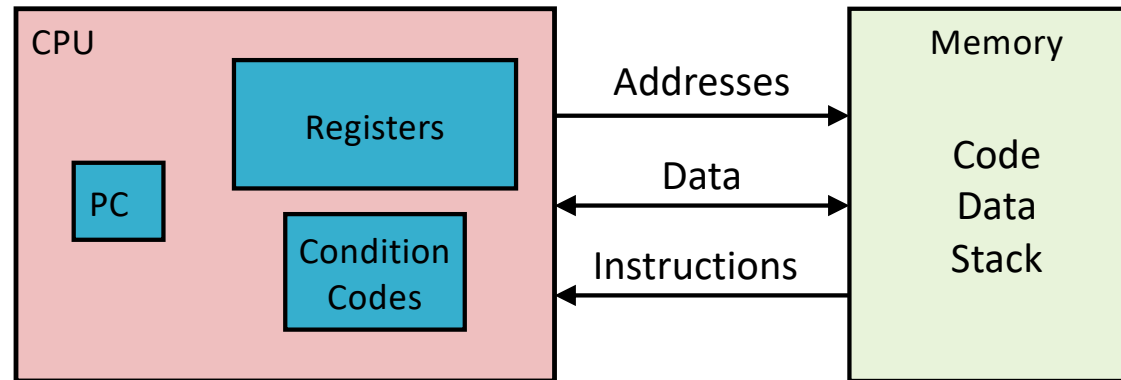


Definitions

- **Architecture:** (also ISA: instruction set architecture) The parts of a processor design that one needs to understand for writing correct machine/assembly code
 - Examples: instruction set specification, registers
 - **Machine Code:** The byte-level programs that a processor executes
 - **Assembly Code:** A text representation of machine code
- **Microarchitecture:** Implementation of the architecture
 - Examples: cache sizes and core frequency
- Example ISAs:
 - Intel: x86, IA32, Itanium, x86-64
 - ARM: Used in almost all mobile phones
 - RISC V: New open-source ISA



Assembly/Machine Code View



Programmer-Visible State

- **PC: Program counter**
 - Address of next instruction
 - Called "RIP" (x86-64)
- **Register file**
 - Heavily used program data
- **Condition codes**
 - Store status information about most recent arithmetic or logical operation
 - Used for conditional branching
- **Memory**
 - Byte addressable array
 - Code and user data
 - Stack to support procedures




Assembly Characteristics: Data Types

- “Integer” data of 1, 2, 4, or 8 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- (SIMD vector data types of 8, 16, 32 or 64 bytes)
- Code: Byte sequences encoding series of instructions
- No aggregate types such as arrays or structures
 - Just contiguously allocated bytes in memory



Next segment Processor registers





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.3. Registers

EL3011 Arsitektur Sistem Komputer



x86-64 Integer Registers

%rax	%eax	%r8	%r8d
%rbx	%ebx	%r9	%r9d
%rcx	%ecx	%r10	%r10d
%rdx	%edx	%r11	%r11d
%rsi	%esi	%r12	%r12d
%rdi	%edi	%r13	%r13d
%rsp	%esp	%r14	%r14d
%rbp	%ebp	%r15	%r15d

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)
- Not part of memory (or cache)



Some History: IA32 Registers


general purpose	%eax	%ax	%ah	%al	accumulate
	%ecx	%cx	%ch	%cl	counter
	%edx	%dx	%dh	%dl	data
	%ebx	%bx	%bh	%bl	base
	%esi	%si			source index
	%edi	%di			destination index
	%esp	%sp			stack pointer
	%ebp	%bp			base pointer

16-bit virtual registers
(backwards compatibility)



Next segment Processor operations





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.4. Processor Operations

EL3011 Arsitektur Sistem Komputer



Assembly Characteristics: Operations

- Transfer data between memory and register
 - Load data from memory into register
 - Store register data into memory
- Perform arithmetic function on register or memory data
- Transfer control
 - Unconditional jumps to/from procedures
 - Conditional branches
 - Indirect branches



Moving Data

- Moving Data

movq Source, Dest

- Operand Types

- **Immediate:** Constant integer data

- Example: \$0x400, \$-533
 - Like C constant, but prefixed with '\$'
 - Encoded with 1, 2, or 4 bytes

- **Register:** One of 16 integer registers

- Example: %rax, %r13
 - But %rsp reserved for special use
 - Others have special uses for particular instructions

- **Memory:** 8 consecutive bytes of memory at address given by register

- Simplest example: (%rax)
 - Various other “addressing modes”

%rax
%rcx
%rdx
%rbx
%rsi
%rdi
%rsp
%rbp
%rN



movq Operand Combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x108, %rax	temp = 0x108;
		Mem	movq \$-147, (%rax)	*p = -147;
	Reg	Reg	movq %rax, %rdx	temp2 = temp1;
		Mem	movq %rax, (%rdx)	*p = temp;
	Mem	Reg	movq (%rax), %rdx	temp = *p;

Cannot do memory-memory transfer with a single instruction



movq Operand Combinations

```
movq $0x108,%rax
```

```
movq $-147, (%rax)
```

```
movq %rax,%rdx
```

```
movq %rax, (%rcx)
```

```
movq (%rax), %rdx
```

Register

%rax
%rcx 0x120
%rdx
%rbx
%rsi
%rdi
%rsp
%rbp
%rN


Memory

Address
0x120
0x118
0x110
0x108
0x100



Next segment more **mov** operations





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.5. mov Operation

EL3011 Arsitektur Sistem Komputer



Simple Memory Addressing Modes

- Normal (R) Mem[Reg[R]]

- Register R specifies memory address
- Aha! Pointer dereferencing in C

```
movq (%rcx) , %rax
```

- Displacement D(R) Mem[Reg[R]+D]

- Register R specifies start of memory region
- Constant displacement D specifies offset

```
movq 8(%rbp) , %rdx
```



Example of Simple Addressing Modes

```
void whatAmI(<type> a, <type> b)
{
    ???
}
```

%rdi

%rsi

```
whatAmI:
    movq    (%rdi), %rax
    movq    (%rsi), %rdx
    movq    %rdx, (%rdi)
    movq    %rax, (%rsi)
    ret
```



Example of Simple Addressing Modes


```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    movq    (%rdi), %rax
    movq    (%rsi), %rdx
    movq    %rdx, (%rdi)
    movq    %rax, (%rsi)
    ret
```



Understanding swap ()

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register	Value
%rdi	xp
%rsi	yp
%rax	t0
%rdx	t1

swap:

```
movq    (%rdi), %rax    # t0 = *xp
movq    (%rsi), %rdx    # t1 = *yp
movq    %rdx, (%rdi)    # *xp = t1
movq    %rax, (%rsi)    # *yp = t0
ret
```



Understanding swap ()

Registers

%rdi	0x120
%rsi	0x100
%rax	
%rdx	

Memory


Address	
0x120	123
0x118	
0x110	
0x108	
0x100	456

```

swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
  
```




Understanding swap ()



```
swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```



Understanding swap ()




```

swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
  
```




Understanding swap ()



```
swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```




Understanding swap ()



```
swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
```





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.6. Addressing Modes

EL3011 Arsitektur Sistem Komputer



Simple Memory Addressing Modes

- Normal (R) Mem[Reg[R]]

- Register R specifies memory address
- Aha! Pointer dereferencing in C

```
movq (%rcx) , %rax
```

- Displacement D(R) Mem[Reg[R]+D]

- Register R specifies start of memory region
- Constant displacement D specifies offset

```
movq 8(%rbp) , %rdx
```



Complete Memory Addressing Modes

- Most General Form

$D(Rb, Ri, S) \quad \text{Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri] + D]$

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for `%rsp`
- S: Scale: 1, 2, 4, or 8



Complete Memory Addressing Modes

- Most General Form

$D(Rb, Ri, S) \quad \text{Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri] + D]$

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for `%rsp`
- S: Scale: 1, 2, 4, or 8 (*why these numbers?*)

- Special Cases

$(Rb, Ri) \quad \text{Mem}[\text{Reg}[Rb] + \text{Reg}[Ri]]$

$D(Rb, Ri) \quad \text{Mem}[\text{Reg}[Rb] + \text{Reg}[Ri] + D]$

$(Rb, Ri, S) \quad \text{Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri]]$



Address Computation Examples

%rdx	0xf000
%rcx	0x0100

D(Rb,Ri,S)

Mem[Reg[Rb]+S*Reg[Ri]+ D]

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for **%rsp**
- S: Scale: 1, 2, 4, or 8 (*why these numbers?*)

Expression	Address Computation	Address
0x8 (%rdx)		
(%rdx, %rcx)		
(%rdx, %rcx, 4)		
0x80 (, %rdx, 2)		



Address Computation Examples

Register


%rdx	0xf000
%rcx	0x0100
%rax	
%rbx	

Memory

	Address
	0xf008
	0xf100
	0xf400
	0x1e080

Expression	Address Computation	Address
0x8(%rdx)	0xf000 + 0x8	0xf008
(%rdx,%rcx)	0xf000 + 0x100	0xf100
(%rdx,%rcx,4)	0xf000 + 4*0x100	0xf400
0x80(,%rdx,2)	2*0xf000 + 0x80	0x1e080





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.7. Addressing Computation

EL3011 Arsitektur Sistem Komputer



Complete Memory Addressing Modes

- Most General Form

$D(Rb, Ri, S) \quad \text{Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri] + D]$

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for `%rsp`
- S: Scale: 1, 2, 4, or 8



Address Computation Examples

%rdx	0xf000
%rcx	0x0100

D(Rb,Ri,S)

Mem[Reg[Rb]+S*Reg[Ri]+ D]

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for **%rsp**
- S: Scale: 1, 2, 4, or 8 (*why these numbers?*)

Expression	Address Computation	Address
0x8 (%rdx)		
(%rdx, %rcx)		
(%rdx, %rcx, 4)		
0x80 (, %rdx, 2)		



Address Computation Examples

Register

%rdx	0xf000
%rcx	0x0100
%rax	
%rbx	

Memory

	Address
	0xf008
	0xf100
	0xf400
	0x1e080

Expression	Address Computation	Address
0x8(%rdx)	0xf000 + 0x8	0xf008
(%rdx,%rcx)	0xf000 + 0x100	0xf100
(%rdx,%rcx,4)	0xf000 + 4*0x100	0xf400
0x80(,%rdx,2)	2*0xf000 + 0x80	0x1e080



Address Computation Instruction

- **leaq** Src, Dst
 - Src is address mode expression
 - Set Dst to address denoted by expression
- Uses
 - Computing addresses without a memory reference
 - E.g., translation of `p = &x[i];`
 - Computing arithmetic expressions of the form $x + k*y$
 - $k = 1, 2, 4, \text{ or } 8$

- Example

```
long m12(long x)
{
    return x*12;
}
```

Converted to ASM by compiler:

```
leaq (%rdi,%rdi,2), %rax # t = x+2*x
salq $2, %rax           # return t<<2
```



Address Computation Instruction

- Example: consider `%rdx = x`

```
movq 7(%rdx,%rdx,4),%rax
```

```
leaq 7(%rdx,%rdx,4),%rax
```


Register

<code>%rax</code>	
<code>%rcx</code>	
<code>%rdx</code>	
<code>%rbx</code>	

Memory

Address





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

5.8. Arithmetic and Logical Operations

EL3011 Arsitektur Sistem Komputer



Some Arithmetic Operations

- Two Operand Instructions:

Format	Computation	
<code>addq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} + \text{Src}$
<code>subq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} - \text{Src}$
<code>imulq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} * \text{Src}$
<code>shlq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} \ll \text{Src}$ Synonym: <code>salq</code>
<code>sarq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} \gg \text{Src}$ Arithmetic
<code>shrq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} \gg \text{Src}$ Logical
<code>xorq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} \wedge \text{Src}$
<code>andq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} \& \text{Src}$
<code>orq</code>	<code>Src, Dest</code>	$\text{Dest} = \text{Dest} \text{Src}$

- Watch out for argument order! *Src, Dest* (Warning: Intel docs use “op *Dest, Src*”)
- No distinction between signed and unsigned int (why?)



Some Arithmetic Operations

- One Operand Instructions

<code>incq</code>	<code>Dest</code>	$\text{Dest} = \text{Dest} + 1$
<code>decq</code>	<code>Dest</code>	$\text{Dest} = \text{Dest} - 1$
<code>negq</code>	<code>Dest</code>	$\text{Dest} = -\text{Dest}$
<code>notq</code>	<code>Dest</code>	$\text{Dest} = \sim\text{Dest}$

- See book for more instructions

- Depending how you count, there are 2,034 total x86 instructions
- (If you count all addr modes, op widths, flags, it's actually 3,683)



Arithmetic Expression Example

```
long arith
(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

arith:

```
    leaq    (%rdi,%rsi), %rax
    addq    %rdx, %rax
    leaq    (%rsi,%rsi,2), %rdx
    salq    $4, %rdx
    leaq    4(%rdi,%rdx), %rcx
    imulq    %rcx, %rax
    ret
```

Interesting Instructions

- **leaq**: address computation
- **salq**: shift
- **imulq**: multiplication
 - Curious: only used once...



```

1 arith:
2   leaq    (%rdi,%rsi), %rax
3   addq    %rdx, %rax
4   leaq    (%rsi,%rsi,2), %rdx
5   salq    $4, %rdx
6   leaq    4(%rdi,%rdx), %rcx
7   imulq   %rcx, %rax
8   ret

```

```

long arith(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}

```

Register

%rax	
%rcx	
%rdx	z
%rdi	x
%rsi	y




Understanding Arithmetic Expression Example

```
long arith
(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

```
arith:
    leaq    (%rdi,%rsi), %rax    # t1
    addq    %rdx, %rax          # t2
    leaq    (%rsi,%rsi,2), %rdx
    salq    $4, %rdx            # t4
    leaq    4(%rdi,%rdx), %rcx   # t5
    imulq    %rcx, %rax          # rval
    ret
```

Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rdx	Argument z , t4
%rax	t1 , t2 , rval
%rcx	t5





STEI - Institut Teknologi Bandung

Modul 5. Intel X86-64 Bit Machine

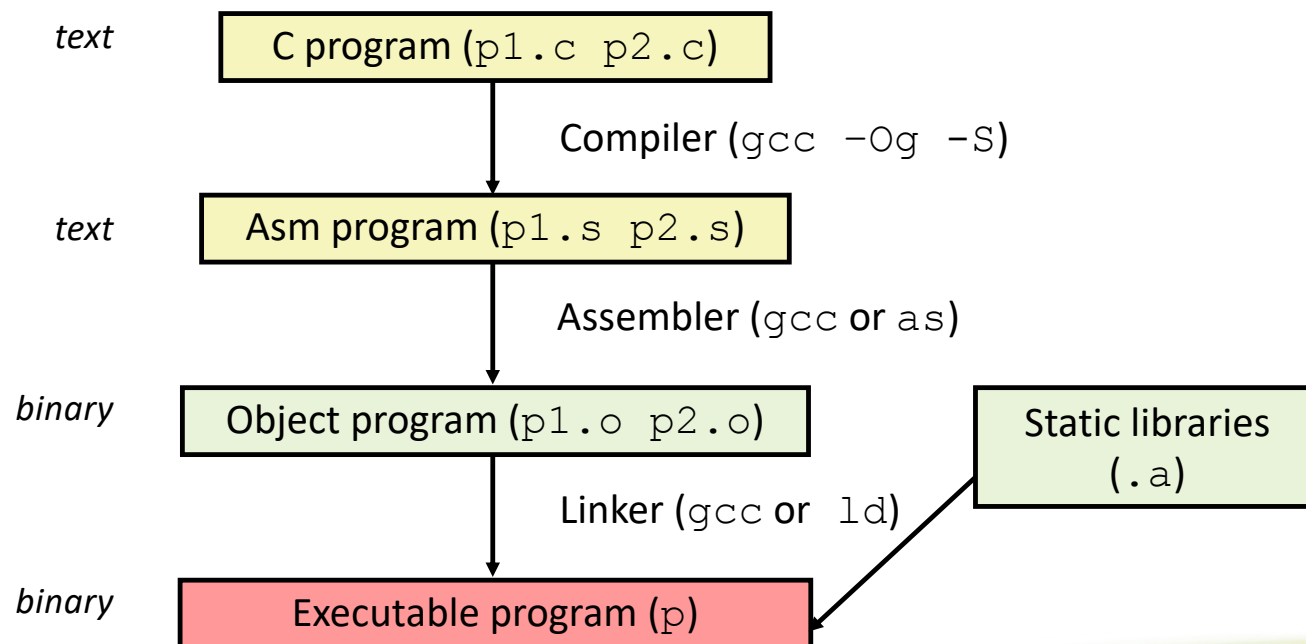
5.9. C Programming, Assembly and Machine Code

EL3011 Arsitektur Sistem Komputer



Turning C into Object Code

- Code in files **p1.c** **p2.c**
- Compile with command: **gcc -Og p1.c p2.c -o p**
 - Use basic optimizations (**-Og**) [New to recent versions of GCC]
 - Put resulting binary in file **p**



Compiling Into Assembly

C Code (swap.c)

```
1 void swap (long *xp, long *yp)
2 {
3     long t0 = *xp;
4     long t1 = *yp;
5     *xp = t1;
6     *yp = t0;
7 }
8
```

Obtain with command

```
gcc -O2 -S swap.c
```

Produces file swap.s

Warning: Will get very different results on other machines due to different versions of gcc and different compiler settings.

Generated x86-64 Assembly

```
Command Prompt

C:\Codes>gcc -O2 -S swap.c

C:\Codes>type swap.s
.file "swap.c"
.text
.p2align 4,,15
.globl swap
.def swap; .scl 2; .type 32; .endef
.seh_proc swap

swap:
.seh_endprologue
movl (%rcx), %eax
movl (%rdx), %r8d
movl %r8d, (%rcx)
movl %eax, (%rdx)
ret
.seh_endproc
.ident "GCC: (x86_64-posix-seh-rev0, Built by MinGW-W64 p

C:\Codes>
```



Compiling Into Assembly

C Code (arith.c)

```
1 long arith(long x, long y, long z)
2 {
3     long t1 = x + y;
4     long t2 = x + t1;
5     long t3 = y + 4;
6     long t4 = y * 48;
7     long t5 = t3 + t4;
8     long rval = t2 * t5;
9     return rval;
10 }
11
```

Generated x86-64 Assembly

```
arith:
    .seh_endprologue
    leal    (%rdx,%rdx,2), %eax
    sall    $4, %eax
    leal    4(%rdx,%rax), %eax
    leal    (%rdx,%rcx,2), %edx
    imull    %edx, %eax
    ret
```



Summary

- History of Intel processors and architectures
 - Evolutionary design leads to many quirks and artifacts
- C, assembly, machine code
 - New forms of visible state: program counter, registers, ...
 - Compiler must transform statements, expressions, procedures into low-level instruction sequences
- Assembly Basics: Registers, operands, move
 - The x86-64 move instructions cover wide range of data movement forms
- Arithmetic
 - C compiler will figure out different instruction combinations to carry out computation

